

Lab: Bootstrap Kubernetes Cluster Using Ansible

Introduction:

Kubernetes is one of the most popular **open-source** and **enterprise-ready** container orchestration systems. It's used to automate the deployment, scaling, and management of containerized applications. Manual **Kubernetes installation** is a laborious and error-prone process. However, it can be dramatically simplified by using configuration management tools such as Ansible. In this Lab let's Learn how to deploy a full-function Kubernetes cluster using Ansible with our installation package

Cluster Designing

Our Kubernetes cluster consists of three servers. One of them will be working as Kubernetes **Controller**. The other two are **worker** nodes. All servers are in the same internal network, 192.168.100.0/24. Software components the cluster depends on are **Kubernetes, Etcd, Docker, WeaveNet**

The following is our server inventory

| Hostname | IP Address | Roles |
|-------------------|-----------------|-------------|
| 1. eoc-controller | 192.168.100.150 | Controller |
| 2. eoc-node1 | 192.168.100.151 | Worker node |
| 3. eoc-node2 | 192.168.100.152 | Worker node |

Objectives

- Verifying Communication with Kubernetes
- Pre-requisites to install the Kubernetes Cluster
- Required Images to work
- Verifying controller status using Ad-Hoc command
- Join token capture
- Adding controller and worker nodes using Ad-Hoc command
- Verifying cluster status using Ad-Hoc command
- Node role

Note: Login to **eoc-controller** as **admin** user with password as **linux**

1. Verifying Communication with Kubernetes

- 1.1 Let's create Ansible inventory file **kube-infra** to tell Ansible how to communicate with the Kubernetes controller and worker nodes.

```
# cat > kube-infra << EOF
[controller]
eoc-controller
[workers]
eoc-node1
eoc-node2
EOF
```

Note: Listing the controller node and the worker nodes in different sections in the hosts file will allow us to target the playbooks at the specific node type later on.

1.2 We can test it's working by doing a Ansible ping.

```
# ansible "controller,workers" -i kube-infra -m ping
```

Output:

```
[admin@eoc-controller ~]$ ansible "controller,workers" -i kube-infra -m ping
eoc-node1 | SUCCESS => {
  "ansible_facts": {
    "discovered_interpreter_python": "/usr/libexec/platform-python"
  },
  "changed": false,
  "ping": "pong"
}
eoc-node2 | SUCCESS => {
  "ansible_facts": {
    "discovered_interpreter_python": "/usr/libexec/platform-python"
  },
  "changed": false,
  "ping": "pong"
}
eoc-controller | SUCCESS => {
  "ansible_facts": {
    "discovered_interpreter_python": "/usr/libexec/platform-python"
  },
  "changed": false,
  "ping": "pong"
}
```

2. Pre-requisites to install the Kubernetes Cluster

2.1 Let's deploy the pre-requisites to install the kubernetes cluster by creating a yaml named **kube-pre-requisites.yml**

Let's target the host and declaring variables

```
1 ---
2 - hosts: "controller, workers"
3   become: yes
4   tasks:
```

2.2 Let's disable the firewalld service.

```
5 - name: disabling firewalld service
6   become: yes
7   service:
8     name: firewalld
9     state: stopped
10    enabled: no
```

2.3 Enable and start the chronyd service.

```
11 - name: Enabling and Starting Chrony
12   become: yes
13   service:
14     name: chronyd
15     state: started
16     enabled: yes
```

2.4 Disabling swap on all nodes.

```
17 - name: Disabling Swap on all nodes
18   become: yes
19   shell: swapoff -a
```

2.5 Commenting out swap functionality completely.

```
20 - name: Commenting Swap entries in /etc/fstab
21   replace:
22     path: /etc/fstab
23     regexp: ' (^/. *swap*) '
24     replace: '# \1'
```

2.6 Enabling br_netfilter.

```
25 - name: Ensure br_netfilter is enabled.
26   modprobe:
27     name: br_netfilter
28     state: present
```

2.7 Updating kernel settings.

```
29 - name: update kernel settings
30   sysctl:
31     name: net.bridge.bridge-nf-call-iptables
32     value: 1
33     sysctl_set: yes
34     state: present
35     reload: yes
```

2.8 Installing and configuring docker.

```
36 - name: "Installing Docker Prerequisite packages"
37   dnf:
38     name:
39       - epel-release
40       - yum-utils
41     state: latest
42 - name: "configuring Yum Repository for docker"
43   become: yes
44   shell: yum-config-manager --add-repo https://download.docker.com/linux/centos/docker-ce.repo
45 - name: " Installing Docker latest version"
46   become: yes
47   shell: dnf install -y docker-ce docker-ce-cli containerd.io --allowprerelease
```

2.9 Configuring crictl.

```

48 - name: Commenting disabled_plugins in /etc/containerd/config.toml
49   replace:
50     path: /etc/containerd/config.toml
51     regexp: '^disabled_plugins'
52     replace: '# disabled_plugins'
53 - name: setting endpoints for crictl
54   become: yes
55   shell: |
56     cat > /etc/crictl.yaml <<EOF
57     runtime-endpoint: unix:///run/containerd/containerd.sock
58     image-endpoint: unix:///run/containerd/containerd.sock
59     EOF

```

2.10 Enabling docker and containerd.

```

60 - name: start and enable containerd service.
61   become: yes
62   service:
63     name: "containerd"
64     state: started
65     enabled: yes
66 - name: "Starting Docker Service"
67   service:
68     name: "docker"
69     state: started
70     enabled: yes

```

2.11 Adding Kubernetes repo and installing.

```

71 - name: Add Kubernetes repository
72   shell: |
73     cat <<EOF | sudo tee /etc/yum.repos.d/kubernetes.repo
74     [kubernetes]
75     name=Kubernetes
76     baseurl=https://pkgs.k8s.io/core:/stable:/v1.28/rpm/
77     enabled=1
78     gpgcheck=1
79     gpgkey=https://pkgs.k8s.io/core:/stable:/v1.28/rpm/repodata/repomd.xml.key
80     EOF
81 - name: "Installing Kubeadm, Kubelet, and Kubectl"
82   dnf:
83     name:
84       - kubectl
85       - kubeadm
86       - kubelet
87     state: present

```

2.12 Starting the kubelet service.

```

88 - name: "Starting Kubelet Service"
89   service:
90     name: "kubelet"
91     state: started
92     enabled: yes

```

2.13 Let's verify the syntax of the yaml kube-pre-requisites.yml.

```
# ansible-playbook -i kube-infra kube-pre-requisites.yml --syntax-check
```

Output:

```
[admin@eoc-controller ~]$ ansible-playbook -i kube-infra kube-pre-requisites.yml --syntax-check
playbook: kube-pre-requisites.yml
```

2.14 Let's run the Playbook kube-pre-requisites.yml.

```
# ansible-playbook -i kube-infra kube-pre-requisites.yml
```

Output:

```
[admin@eoc-controller ~]$ ansible-playbook -i kube-infra kube-pre-requisites.yml

PLAY [controller, workers] *****

TASK [Gathering Facts] *****
ok: [eoc-node1]
ok: [eoc-node2]
ok: [eoc-controller]

TASK [disabling firewalld service] *****
ok: [eoc-node2]
ok: [eoc-controller]
changed: [eoc-node1]

TASK [Enable and start chronyd service] *****
changed: [eoc-node2]
changed: [eoc-node1]
changed: [eoc-controller]

TASK [Disabling Swap on all nodes] *****
changed: [eoc-node1]
changed: [eoc-node2]
changed: [eoc-controller]

TASK [Commenting Swap entries in /etc/fstab] *****
changed: [eoc-node2]
changed: [eoc-node1]
changed: [eoc-controller]

.....

TASK [Installing Kubeadm, Kubelet, and Kubectl] *****
ok: [eoc-controller]
changed: [eoc-node2]
changed: [eoc-node1]

TASK [Starting Kubelet Service] *****
ok: [eoc-controller]
changed: [eoc-node1]
changed: [eoc-node2]

PLAY RECAP *****
eoc-controller      : ok=18   changed=7   unreachable=0   failed=0   skipped=0   rescued=0   ignored=0
eoc-node1           : ok=18   changed=9   unreachable=0   failed=0   skipped=0   rescued=0   ignored=0
eoc-node2           : ok=18   changed=9   unreachable=0   failed=0   skipped=0   rescued=0   ignored=0
```

3. Required Images to work**3.1 Add the ansible-galaxy collection of Kubernetes module if not present.**

```
# ansible-galaxy collection list | grep kubernetes
```

Output:

```
[admin@eoc-controller ~]$ ansible-galaxy collection list | grep kubernetes
kubernetes.core          3.0.0
kubernetes.core          2.3.2
```

Info: If not present use this command to install or download the latest one.

```
# ansible-galaxy collection install kubernetes.core --force
```

Output:

```
[admin@eoc-controller ~]$ ansible-galaxy collection install kubernetes.core --force
Starting galaxy collection install process
Process install dependency map
Starting collection install process
Downloading https://galaxy.ansible.com/api/v3/plugin/ansible/content/published/collections/artifacts/kubernetes-core-3.0.0.tar.gz to /home/admin/.ansible/tmp/ansible-local-24398xeeo7op9/tmpilflmq5/kubernetes-core-3.0.0-wlfpb00
Installing 'kubernetes.core:3.0.0' to '/home/admin/.ansible/collections/ansible_collections/kubernetes/core'
kubernetes.core:3.0.0 was installed successfully
```

3.2 Let's pull the images requires to work as a control-plane on eoc-node1 by creating a yaml named **bootstrapping-controller.yml**.

```
1 ---
2 - hosts: controller
3   become: yes
4   vars:
5     token_file: join_token
```

3.3 Let set the task if something goes wrong it will reset the kubeadm.

```
6   tasks:
7     - name: Resetting kubeadm
8       shell: kubeadm reset -f
9       register: output
```

3.4 Let initialize the kubernetes cluster.

```
10    - name: Initialize the deployment of Kubernetes cluster
11      shell: kubeadm init
12      register: output
```

3.5 Setting up the configuration files.

```
13    - name: "Configuration Files Setup"
14      file:
15        path: "$HOME/.kube"
16        state: directory
17    - name: "Copying Configuration File"
18      copy:
19        src: /etc/kubernetes/admin.conf
20        dest: $HOME/.kube/config
21        remote_src: yes
```

3.6 Changing kubeconfig file permissions.

```
22    - name: Change kubeconfig file permission
23      file:
24        path: $HOME/.kube/config
25        owner: "{{ ansible_effective_user_id }}"
26        group: "{{ ansible_effective_group_id }}"
```

3.7 Let's verify the syntax of the yaml **bootstrapping-controller.yml**

```
# ansible-playbook -i kube-infra bootstrapping-controller.yml
--syntax-check
```

Output:

```
[admin@eoc-controller ~]$ ansible-playbook -i kube-infra bootstrapping-master.yml --syntax-check
playbook: bootstrapping-master.yml
```

3.8 Let's run the Playbook bootstrapping-controller.yml.

```
# ansible-playbook -i kube-infra bootstrapping-controller.yml
```

Output:

```
[admin@eoc-controller ~]$ ansible-playbook -i kube-infra bootstrapping-controller.yml

PLAY [controller] *****
TASK [Gathering Facts] *****
ok: [eoc-controller]

TASK [Resetting kubeadm] *****
changed: [eoc-controller]

TASK [Initialize the deployment of Kubernetes cluster] *****
changed: [eoc-controller]

TASK [Configuration Files Setup] *****
changed: [eoc-controller]

TASK [Copying Configuration File] *****
changed: [eoc-controller]

TASK [Change kubeconfig file permission] *****
ok: [eoc-controller]

TASK [Downloading CNI Plugin] *****
changed: [eoc-controller]

TASK [Storing Logs and Generated token for future purpose.] *****
changed: [eoc-controller -> localhost]

PLAY RECAP *****
eoc-controller      : ok=8    changed=6    unreachable=0    failed=0    skipped=0    rescued=0    ignored=0
```

3.9 Installing the cni plugin.

```
# ansible controller -i kube-infra -m command -a 'kubectl
apply -f
https://github.com/weaveworks/weave/releases/download/v2.8.1
/weave-daemonset-k8s-1.11.yaml'
```

4. Verifying the cluster.**4.1 Let's verify the controller status using ad-hoc command.**

```
# ansible controller -i kube-infra -m command -a 'kubectl
get nodes'
```

Output:

```
[admin@eoc-controller ~]$ ansible controller -i kube-infra -m command -a 'kubectl get nodes'
eoc-controller | CHANGED | rc=0 >>
NAME           STATUS    ROLES          AGE    VERSION
eoc-controller Ready     control-plane  41s    v1.28.6
```

4.2 Get the join token from the output saved from the above playbook.

```
# cat join_token
```


Output:

```
You should now deploy a pod network to the cluster.
Run "kubectl apply -f [podnetwork].yaml" with one of the options listed at:
https://kubernetes.io/docs/concepts/cluster-administration/addons/

Then you can join any number of worker nodes by running the following on each as root:

kubeadm join 192.168.100.150:6443 --token e9f98y.np8w4pq01z4p2d7o \
--discovery-token-ca-cert-hash sha256:cf41c09a6765fe2b072853dc7b6943e55291c96e1e7c42b
```

4.3 Let's join the worker nodes to the controller using ad-hoc command.

```
# ansible workers -i kube-infra -m command -a 'kubeadm join
192.168.100.150:6443 --token 6ms418.ytixb4sn41v7d6b8 --
discovery-token-ca-cert-hash
sha256:b44451fd3211c0b2615a87a1738c5e23c1b498a9bea697208138f
ca2a5c8728e'
```

Output:

```
eoc-node2 | CHANGED | rc=0 >>
[preflight] Running pre-flight checks
[preflight] Reading configuration from the cluster...
[preflight] FYI: You can look at this config file with 'kubectl -n kube-system get cm kubeadm-config -o
[kubelet-start] Writing kubelet configuration to file "/var/lib/kubelet/config.yaml"
[kubelet-start] Writing kubelet environment file with flags to file "/var/lib/kubelet/kubeadm-flags.env"
[kubelet-start] Starting the kubelet
[kubelet-start] Waiting for the kubelet to perform the TLS Bootstrap...

This node has joined the cluster:
* Certificate signing request was sent to apiserver and a response was received.
* The Kubelet was informed of the new secure connection details.

Run 'kubectl get nodes' on the control-plane to see this node join the cluster.
eoc-node1 | CHANGED | rc=0 >>
[preflight] Running pre-flight checks
[preflight] Reading configuration from the cluster...
[preflight] FYI: You can look at this config file with 'kubectl -n kube-system get cm kubeadm-config -o
[kubelet-start] Writing kubelet configuration to file "/var/lib/kubelet/config.yaml"
[kubelet-start] Writing kubelet environment file with flags to file "/var/lib/kubelet/kubeadm-flags.env"
[kubelet-start] Starting the kubelet
[kubelet-start] Waiting for the kubelet to perform the TLS Bootstrap...

This node has joined the cluster:
* Certificate signing request was sent to apiserver and a response was received.
* The Kubelet was informed of the new secure connection details.

Run 'kubectl get nodes' on the control-plane to see this node join the cluster.
```

5. Verifying cluster status using Ad-Hoc command**5.1 Let's verify the Cluster status using ad-hoc command.**

```
# ansible controller -i kube-infra -m command -a 'kubectl
get nodes'
```

Output:

```
[admin@eoc-controller ~]$ ansible controller -i kube-infra -m command -a 'kubectl get nodes'
eoc-controller | CHANGED | rc=0 >>
NAME             STATUS    ROLES    AGE     VERSION
eoc-controller    Ready     control-plane   3m3s    v1.28.6
eoc-node1         Ready     <none>        25s     v1.28.6
eoc-node2         Ready     <none>        25s     v1.28.6
```

5.2 Let's label the nodes by executing below command.

```
# ansible controller -i kube-infra -m command -a 'kubectl
label node eoc-node1 node-role.kubernetes.io/node=''
```



```
# ansible controller -i kube-infra -m command -a 'kubectl
label node eoc-node2 node-role.kubernetes.io/node='
```

Output:

```
[admin@eoc-controller ~]$ ansible controller -i kube-infra -m command -a 'kubectl label node eoc-node1 n
ode-role.kubernetes.io/node='
eoc-controller | CHANGED | rc=0 >>
node/eoc-node1 labeled
[admin@eoc-controller ~]$ ansible controller -i kube-infra -m command -a 'kubectl label node eoc-node2 n
ode-role.kubernetes.io/node='
eoc-controller | CHANGED | rc=0 >>
node/eoc-node2 labeled
```

5.3 Let's verify the Cluster status using ad-hoc command

```
# ansible controller -i kube-infra -m command -a 'kubectl
get nodes'
```

Output:

```
[admin@eoc-controller ~]$ ansible controller -i kube-infra -m command -a 'kubectl get nodes'
eoc-controller | CHANGED | rc=0 >>
NAME                STATUS    ROLES    AGE     VERSION
eoc-controller       Ready     control-plane   4m46s   v1.28.6
eoc-node1            Ready     node           2m8s    v1.28.6
eoc-node2            Ready     node           2m8s    v1.28.6
```